

# Model Selection through Active Learning: Do people use simple heuristics or integrative decision strategies?

Paula Parpart, Eric Schulz  
Maarten Speekenbrink & Bradley C. Love

University College London, Department of Experimental Psychology

June 22, 2015


Questions in Behavioural Science

# If humans are active learners, then...

- 1 ...how can they learn a heuristic that is made up of building blocks on the fly when the goal is to make good decisions?






- 2 ...how can we use that fact to test different decision making strategies in an active learning context?



# Overview

## 1 Distinguishing active versions of decision strategies

Collaborators:  Parpart  Speekenbrink  Love

- How to distinguish between active versions of strategies?
- How does the used strategy depend on the environment?

# PART 1: Can heuristics emerge/grow adaptively?



- Heuristics are made of smaller building blocks
- Different combinations of blocks produce the heuristic toolbox
- Definition of a heuristic:

*"Ignores information to be faster and/or more accurate."*

*"Exhibits a starting rule, a search rule, and a stopping rule."*

## PART 2: Can we distinguish different active algorithms?



- Agent has evolved to obey a certain decision algorithm
- Agent learns with the goal to apply that algorithm
- Definition of active learning:

*“Stepwise selection of observations in order to learn faster.”*

*“Active learning leads to a banana shaped learning curve.”*

## PART 2: Problem statement

### Problem

Given a number of binary cues  $C_1, C_2, \dots, C_n$  –some of which might be invalid– learn how to predict a binary outcome  $y$  by selecting observations as wisely as possible.

### Current heuristics cannot deal with that, because...

- 1 it is not clear how cue validities are learned
- 2 it is not clear how information is ignored
- 3 there are no active versions of them
- 4 active learning has never been utilized for model comparison


## PART 2: Model 1: Active TTB

- **Imagine a scenario without noise:**

Every observation is deterministic and your only goal is, given your hypothesis space over all cue orders,

$\mathcal{H} = \{\text{cue-order}_1, \text{cue-order}_2, \dots, \text{cue-order}_k\}$ , to select the observation  $s^* = \operatorname{argmax}\{|\mathcal{H}| - |\mathcal{H}|s|\}$

- Unfortunately, observations never come without noise
- We have to find a probabilistic version of the same algorithm



## PART 2: Model 1: Active TTB algorithm

- Put a pseudo-count  $\pi$  over all possible cue orders
- If a cue order makes a correct prediction, then  $\pi + +$
- Calculate current entropy over all cue orders  $S_0 = \sum_i p_i \log p_i$
- For every possible comparison  $s$  calculate  $p(y = w)$  and  $p(y = L)$  by  $\pi$ -weighted sum over all cue orders
- For every comparison  $s$ , calculate posterior expected entropy  $\mathbb{E}[S|s] = S(y = W) \times p(y = W) + S(y = L) \times p(y = L)$
- Choose  $s^* = \operatorname{argmax}\{S_0 - \mathbb{E}[S|s]\}$ , that is the observation with the highest expected information gain
- Method works well a priori



## PART 2: Model 2: Active Logistic Regression


- Given a Bayesian variant of logistic regression:

$$f(x) = \frac{1}{1 + \exp(-(\beta_0 + \sum_k \beta_k x_k))}$$

- Calculate the current sum of coefficients' uncertainty  
 $S = \sum_k \mathbb{V}(\beta_k)$
- For every comparison  $s$ , calculate  $p(y = w|s)$  and  $p(y = l|s)$
- For every comparison and outcome calculate  $S|s, y = w$  and  $S|s, y = l$ , that is the reduction of variance
- For every comparison  $s$ , calculate posterior expected uncertainty  
 $\mathbb{E}[S|s] = S(y = W) \times p(y = W) + S(y = L) \times p(y = L)$
- Choose  $s^* = \operatorname{argmax}\{S_0 - \mathbb{E}[S|s]\}$ , that is the observation with the highest expected uncertainty reduction
- Method works well a priori


## Part 2: Alien Olympics: Design

- 1 Learn how well aliens perform in Olympics
- 2 Aliens vary on 4 different features:  
Wings, Camouflage, Diamonds, and Antennæ
- 3 30 learning trials: select 2 out of 4 aliens to compete against each other (\$0.5 reward)
- 4 10 test trials: select 1 out of 2 aliens for your team (bonus dependent on team)
- 5 Environment set up with different levels of compensatoriness





## Part 2: Alien Olympics: Results in test set

- 1 Underlying environment obeys logistic regression
- 2 Weights are generated by a stick-breaking process  
 $\beta'_k \sim \text{Beta}(1, \theta)$   
Define  $\{\beta'_k\}_{k=1}^4$  as:  
 $\beta_k = \beta'_k \prod_{i=1}^{k-1} (1 - \beta'_i)$
- 3 Allows to trade-off different levels of compensatoriness





## Part 2: Alien Olympics: Results of passive part



- 1 Non-compensatory conditions produce lower average score
- 2 Hard to distinguish between models based on test set alone

## Part 2: Alien Olympics: Results of active part



- 1 Participants perform simple but sensible queries
- 2 Can distinguish between active models:  
Only active TTB seems to match behavior well

## PART 2: Conclusion



Active learning can be used to distinguish different models



It is possible to design active versions of classic models



In a first experiment, active TTB seems to do best



Simpler follow-up, Exploitation scenarios, Non-parametrics

# Overall conclusion

## **I hope to have convinced you that...**

- 1 ...active learning is an exciting tool to add to our methodological repertoire
- 2 ...it is both possible and fruitful to design active learning algorithms based on classic models
- 3 ...if we want to make claims about cognitive strategies, we also have to make claims about how those are acquired

## **Future steps could involve...**

- 1 ...coming up with more active learning models
- 2 ...designing additional active learning algorithms
- 3 ...modeling both exploration and exploitation scenarios

# Thank you!

## 1 Collaborators:



Schulz



Speekenbrink



Love